

Estimations

S20426

2025-03-15

```
# Generate synthetic data from an exponential distribution
set.seed(123)
data <- rexp(100, rate = 0.1)

# Likelihood function for the exponential distribution
log_likelihood <- function(lambda) {
  sum(dexp(data, rate = lambda, log = TRUE))
}

# Maximum Likelihood Estimation using optim
result <- optim(par = 0.1, log_likelihood, method = "Brent", lower = 0.001, upper = 10)
mle_lambda <- result$par

# Print the MLE estimate
cat("Maximum Likelihood Estimate for Lambda:", mle_lambda, "\n")
```

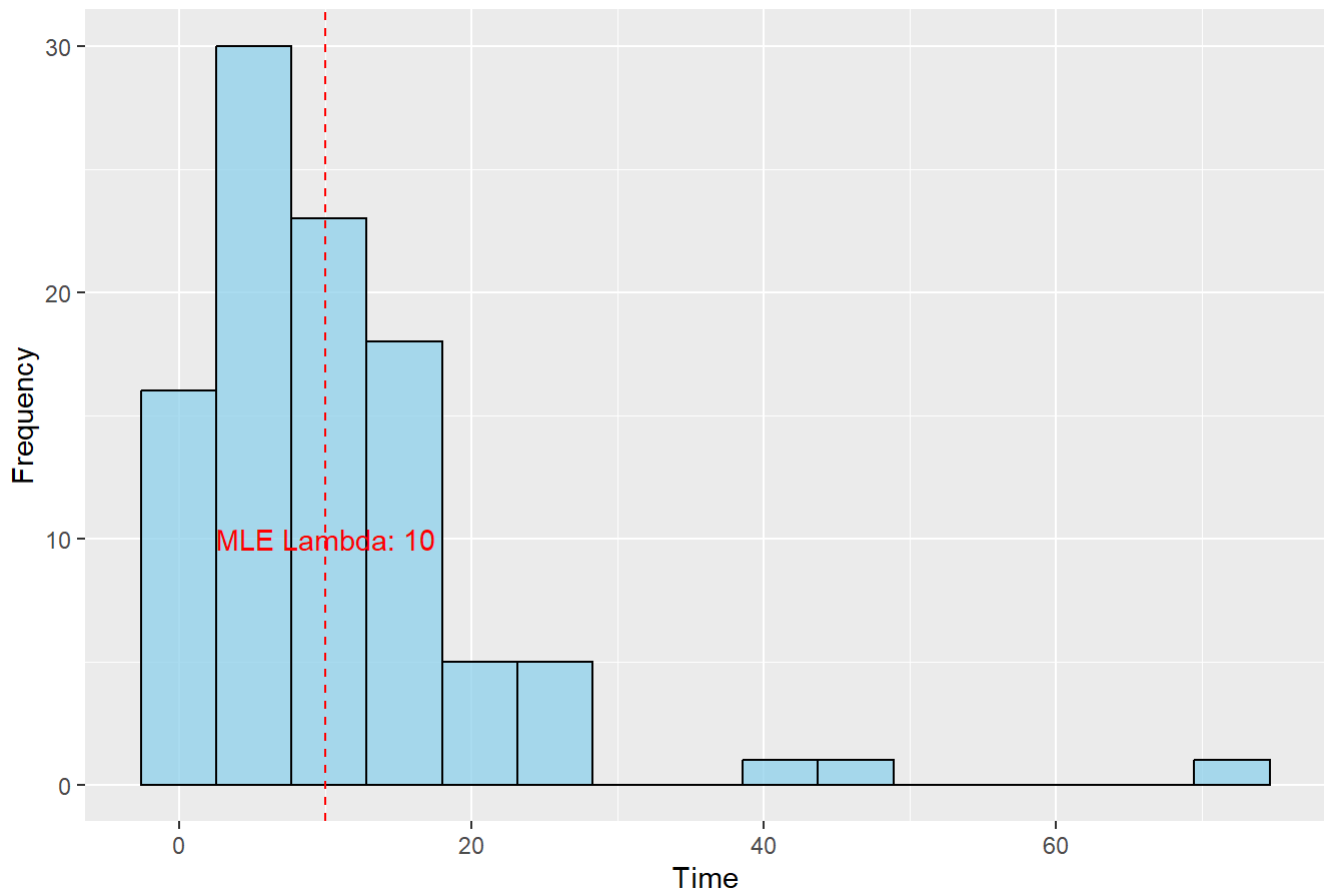
```
## Maximum Likelihood Estimate for Lambda: 10
```

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 4.3.3
```

```
# Plot the histogram of the data with the estimated Lambda
ggplot(data = NULL, aes(x = data)) +
  geom_histogram(bins = 15, fill = "skyblue", color = "black", alpha = 0.7) +
  geom_vline(xintercept = mle_lambda, color = "red", linetype = "dashed") +
  annotate("text", x = mle_lambda + 0.1, y = 10,
    label = paste("MLE Lambda:", round(mle_lambda, 2)), color = "red") +
  labs(title = "Histogram with Estimated Lambda", x = "Time", y = "Frequency")
```

Histogram with Estimated Lambda



Example 01

```
# ?optimize

log_likelihood <- function(p) {

  x <- 103
  n <- 200

  if (p == 0 || p == 1) {
    return(-Inf) # To avoid log(0) which is undefined
  }
  return(x * log(p) + (n - x) * log(1 - p))
}

result <- optimize(log_likelihood, interval = c(0, 1), maximum = TRUE)
result
```

```
## $maximum
## [1] 0.5150064
##
## $objective
## [1] -138.5394
```

Example 02

```
weights <- c(59.001, 38.267, 41.025, 35.555, 46.690, 20.994, 39.407, 52.780,
            57.495, 52.416, 60.062, 48.149, 40.182, 50.929, 49.472, 49.197,
            43.459, 40.493, 60.196, 58.590, 53.645, 53.837, 61.134, 62.115,
            46.517, 41.404, 56.500, 53.281, 44.821, 47.610, 51.178, 58.315,
            34.411, 47.795, 41.828, 60.767, 60.797, 51.421, 51.570, 48.313,
            47.310, 58.078, 38.753, 35.692, 50.604, 42.070, 53.403, 47.405,
            36.952, 53.682)

normal <- function(x,data){

  return (
    sum(dnorm(x,mean(data),sd(data),log = TRUE))
  )
}
result <- optimize(normal,interval = c(-100,+100),data=weights, maximum = TRUE)
result
```

```
## $maximum
## [1] 48.71134
##
## $objective
## [1] -3.086087
```

Confidence Interval for Mean

Case 1: When data is normal/ large sample and σ is known.

```
# set.seed(42) Ensuring reproducibility in simulations.
#same sequence of random numbers is generated every time you run the code.

set.seed(20)
sample_size<-500
pop_div<-10
weight<-sample(45:80,size=sample_size,replace = T)
sample_mean<- mean(weight)
z_critical <- qnorm(0.975) # calculate the z - critical value
margin_error <-z_critical*(pop_div/sqrt(sample_size))

vec<-c(sample_mean - margin_error,sample_mean+ margin_error)
vec
```

```
## [1] 62.28748 64.04052
```

Case 2 : When data is normal / large sample and σ is unknown

```

set.seed(20)
large_sample_weight <- sample(weight,150)
large_sample_t_critical <- qt(0.975,df=149) # find the t value
large_sample_mean <- mean(large_sample_weight)
large_sample_stdev <- sd(large_sample_weight)

large_sample_margin_of_error <- large_sample_t_critical*(large_sample_stdev/sqrt(150))

large_sample_confidence_interval <- c(large_sample_mean- large_sample_margin_of_error,large_s
ample_mean + large_sample_margin_of_error)

large_sample_confidence_interval

```

```
## [1] 61.32493 64.87507
```

Case 3 : When data is non-normal /small samples

#When the data is non-normal or when sample sizes are small, traditional parametric methods (like #using normal distribution assumptions) may not be reliable. Bootstrapping is a resampling technique #that helps estimate the sampling distribution of a statistic (e.g., the mean) without assuming #normality.

```

library(boot)

blood_pressure <- c(72,66,64,66,40,74,50,70,96,92,74,80,60,72,84,74,80,88,94)

mean_fn <- function(x ,indices){
  return (mean(x[indices]))
}
#Performs bootstrap resampling 999 times (R=999).
#The argument indices represents the randomly selected indices in each bootstrap resample.

level.boot <- boot(blood_pressure , mean_fn ,R=999)
boot.ci(level.boot,conf = 0.95)

```

```
## Warning in boot.ci(level.boot, conf = 0.95): bootstrap variances needed for
## studentized intervals
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 999 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = level.boot, conf = 0.95)
##
## Intervals :
## Level      Normal      Basic
## 95%   (67.32, 79.87 )   (67.58, 80.00 )
##
## Level      Percentile      BCa
## 95%   (66.95, 79.37 )   (66.95, 79.37 )
## Calculations and Intervals on Original Scale
```

Confidence Interval for Difference of Means

Case 1: Sampling from two independent normal distributions with known variances.

`x` —> A numeric vector representing the sample data for a one-sample test or the first sample in a two-sample test.

`y` —> (Optional) A numeric vector for the second sample in a two-sample Z-test. If NULL, a one-sample Z-test is performed.

`alternative` —> Specifies the alternative hypothesis. Options: "two.sided" (default), "less" (one-tailed, left-side test), or "greater" (one-tailed, right-side test).

`sigma.x` —> The known population standard deviation for sample `x`. If NULL, an estimate from the sample is used, but this is not a true Z-test.

`sigma.y` —> The known population standard deviation for sample `y` (if performing a two-sample test). If NULL, an estimate from `y` is used.

`mu` —> The hypothesized population mean for a one-sample test (default is 0). `conf.level` The confidence level for the confidence interval (default is 0.95 or 95%).

```
#library("BSDA")
sample1 <- c(52, 55, 49, 50, 53)
sample2 <- c(47, 50, 48, 51, 49)

# Perform two-sample Z-test
z.test(x = sample1, y = sample2, sigma.x = 3, sigma.y = 2, alternative = "two.sided")
```

```
##
## Two-sample z-Test
##
## data: sample1 and sample2
## z = 1.7365, p-value = 0.08248
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -0.360347 5.960347
## sample estimates:
## mean of x mean of y
## 51.8 49.0
```

Case 2: Sampling from two independent normal distributions with unknown variances (small samples).

when population variances are equal

```
sample1 <- c(52, 55, 49, 50, 53)
sample2 <- c(47, 50, 48, 51, 49)

# Perform two-sample Z-test
t.test(x = sample1, y = sample2, alternative = "two.sided", var.equal=TRUE)
```

```
##
## Two Sample t-test
##
## data: sample1 and sample2
## t = 2.1864, df = 8, p-value = 0.06026
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -0.1531262 5.7531262
## sample estimates:
## mean of x mean of y
##      51.8      49.0
```

when population variances are unequal

```
sample1 <- c(52, 55, 49, 50, 53)
sample2 <- c(47, 50, 48, 51, 49)

# Perform two-sample Z-test
? t.test()
```

```
## starting httpd help server ... done
```

```
t.test(x = sample1, y = sample2, alternative = "two.sided", var.equal=FALSE)
```

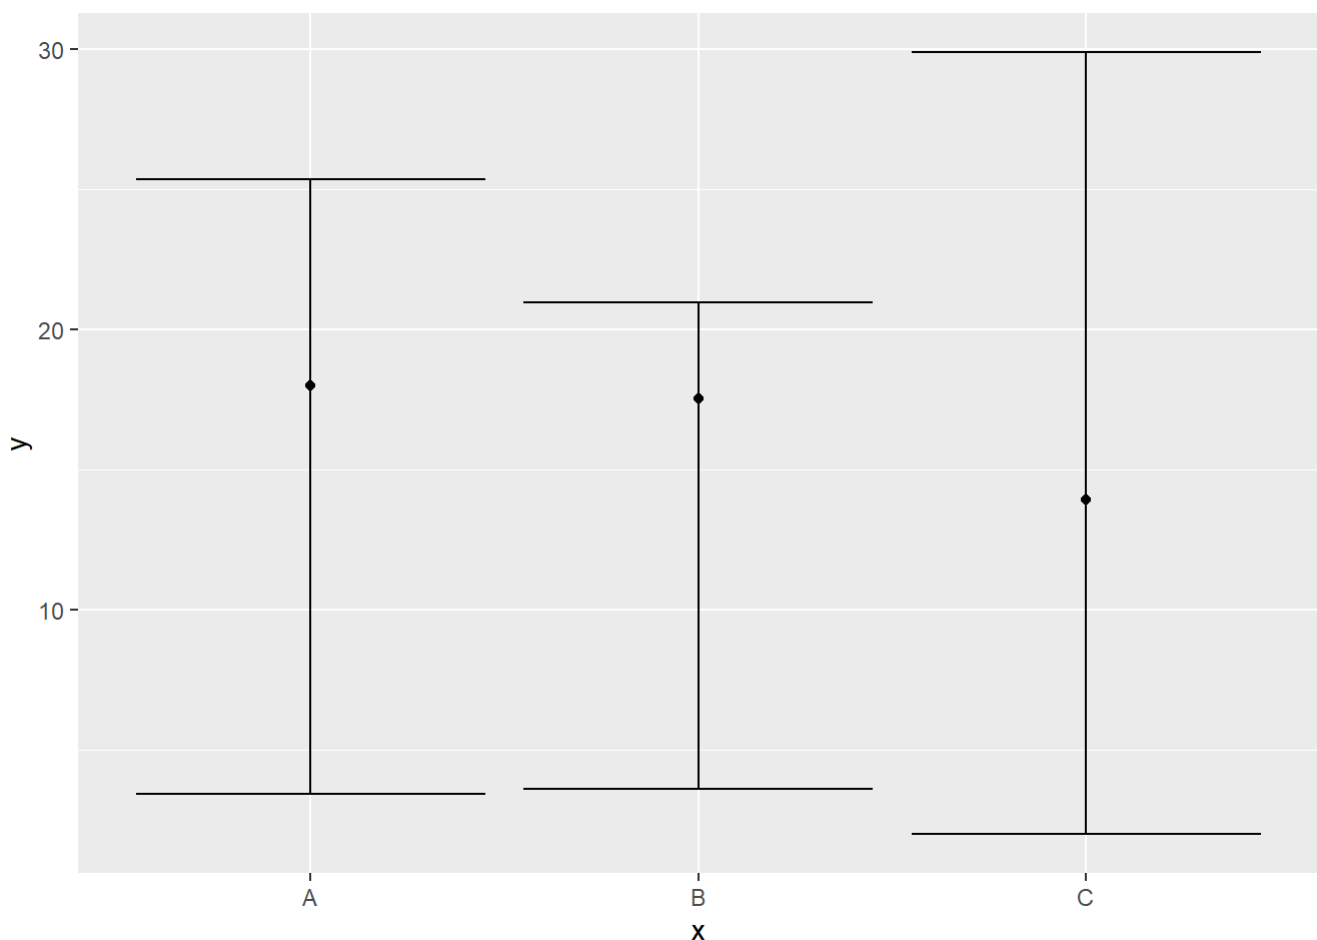
```
##
## Welch Two Sample t-test
##
## data: sample1 and sample2
## t = 2.1864, df = 6.9427, p-value = 0.06534
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -0.2332708 5.8332708
## sample estimates:
## mean of x mean of y
##      51.8      49.0
```

```

set.seed(123456)                                # Create example data
data <- data.frame(x = c("A","B","C"),
                  y = round(runif(3, 10, 20),2),
                  lower = round(runif(3, 0, 10),2),
                  upper = round(runif(3, 20, 30),2))

library(ggplot2)
ggplot(data, aes(x, y)) +                        # ggplot2 plot with confidence intervals
  geom_point() +
  geom_errorbar(aes(ymin = lower, ymax = upper))

```



Confidence Intervals for Proportion

Case 1: For large sample (Using Normal approximation)

```

set.seed(10)
hair_col <- c(rep("black",1500),rep("brown",1000),rep("blonde",500))
sampleP <- sample(hair_col,1000)
Htable <- table(sampleP)
Htable

```

```

## sampleP
##  black blonde  brown
##   498   176   326

```

```
prop.table(Htable)
```

```
## sampleP
## black blonde brown
## 0.498 0.176 0.326
```

```
z=qnorm(0.975)
p=0.498
n=1000

margin_error = z*sqrt(p*(1-p)/n)
Interval=c(p-margin_error,p+margin_error)
Interval
```

```
## [1] 0.4670105 0.5289895
```

Case 1: For large sample (Using Binomial Distribution)

```
library(epitools)

binom.exact(x=48,n=100,conf.level=0.95)
```

```
##      x    n proportion      lower      upper conf.level
## 1 48 100          0.48 0.3790055 0.5822102          0.95
```

```
binom.wilson(x=498,n=100,conf.level = 0.95)
```

```
## Warning in sqrt(((x * (n - x))/n^3) + Z^2/(4 * n^2)): NaNs produced
```

```
##      x    n proportion lower upper conf.level
## 1 498 100          4.98  NaN  NaN          0.95
```

```
binom.approx(x=498,n=100,conf.level = 0.95)
```

```
## Warning in sqrt(x * (n - x)/(n^3)): NaNs produced
```

```
##      x    n proportion lower upper conf.level
## 1 498 100          4.98  NaN  NaN          0.95
```

Case 2: For small sample (Using Binomial Distribution)

When sample size is small, confidence interval for population can be calculated using `binom.test()` function.

```
gender = c("f","f","f","m","m","f","f","m","m","f")
table(gender)
```

```
## gender
## f m
## 6 4
```



```
binom.test(6,10,conf.level = 0.95)
```

```
##  
## Exact binomial test  
##  
## data: 6 and 10  
## number of successes = 6, number of trials = 10, p-value = 0.7539  
## alternative hypothesis: true probability of success is not equal to 0.5  
## 95 percent confidence interval:  
## 0.2623781 0.8784477  
## sample estimates:  
## probability of success  
## 0.6
```